

Convolutional neural networks (CNNs) on Pneumonia Detection through Chest X-rays

Joy Zeng

2020-12-04

1 Introduction

- Background
- Data

2 Methodology

- Data Preprocessing
- Models
 - Supervised learning framework
 - More details on CNN
- Visualizing the ConvNet via Saliency maps
- Implementation

3 Experiments and Results

4 Conclusion

Introduction

- **Background:**

- Pneumonia is an infection of the lungs
- Diseases such as COVID-19, SARS onset pneumonia in both lungs
- Chest X-rays is often used in detecting lung infection

- **Data:**

- Chest X-ray that are labeled with either healthy or pneumonia
- <https://www.kaggle.com/praveengovi/coronahack-chest-xraydataset>



Figure: An Example Image in the dataset

- **training set:** 5286 images (3944 Pneumonia)
test set: 624 images (390 Pneumonia)
- **Goal:** classify the X-rays into Healthy vs Pneumonia

- 1 **Resized** all of the images to the size of 300×300
 - Since model training requires all the images are of the same size
- 2 **Normalized** the images before modeling (subtract both train/test sets by empirical mean and divided by standard deviation **computed over the training data**)
 - Classification loss is sensitive to the changes in weight matrix
- 3 Split training data into $\begin{cases} 50\% \text{ training set} \\ 50\% \text{ validation set} \end{cases}$

Models (Supervised learning framework)

Supervised learning framework:

- Given a dataset: $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$, the goal is to train a predictive function $f_{\theta}(\cdot)$ so that $f_{\theta}(\mathbf{x})$ will be able to predict labels y_i (healthy or Pneumonia).
- Define **Loss** (error measurement) as: $L(\theta) = \frac{1}{N} \sum_i L_i(f(\mathbf{x}_i; \theta), y_i)$ where $L_i = -\log\left(\frac{e^{f_{y_i}}}{\sum_j e^{f_j}}\right)$ (**Cross-entropy loss**/Softmax classifier)
- Optimization:**
We want to find a set of parameters that minimize the loss function:
$$\hat{\theta} = \arg \min_{\theta \in \Theta} \frac{1}{n} \sum_{i=1}^N L(y_i, f_{\theta}(\mathbf{x}_i))$$
 - Most basic algorithm: **stochastic gradient descent**
 - SGD performs poorly on high-dimensional non-convex optimization. so we used the improved version, **Adam**, [Kingma and Ba (2014)]
- Regularization** to prevent overfitting (more details later)
- f here is CNN models, and \mathbf{x} 's are matrix representation of pixels

Models (Backprop)

- The key of SGD algorithm is **computing the gradients** of the loss with respect to the model parameters.

$$L(\theta) = \frac{1}{N} \sum_i L_i(f(\mathbf{x}_i; \theta))$$

$$\nabla_{\theta} L(\theta) = \frac{1}{N} \sum_i \nabla_{\theta} L_i(f(\mathbf{x}_i; \theta), y_i)$$

- Hard to compute gradients in deep networks
- Use **Backpropagation** to compute the gradients
 - recursively applies the chain rule during the backward pass to compute the gradients

Models (More details on CNN)

Architecture of CNN:

A series of convolutional layer + pooling layer, added by fully connected layers at the end.




Figure: The General Architecture of CNN

- 1 Use **ReLU** as our activation function for non-saturating nonlinearity
- 2 Weight Initialization: use **He Init** specifically derived for ReLU, [He et al. (2015)]
- 3 Use **Batch normalization** to improve generalization on the test set (Regularization), [Lofe and Szegedy (2015)]
 - Normalizes each layer with zero-mean and unit variance
 - Robust to bad initialization
- 4 Dropout (Regularization)

Models (More details on CNN)

We explored the followings:

- ① (fully connected layer) \times 1 as the baseline model
- ② (2D convolution layer + dropout + max/average pooling) \times 2 + fully connected layer
- ③ (2D convolution layer + dropout + batch normalization + max/average pooling) \times 2 + (fully connected layer) \times 2

In particular:

- The final classification performance was evaluated based on accuracy, recall, and precision on the test set.

Visualizing the ConvNet via Saliency maps

To understand the networks:

- interested in knowing which pixels in the input image matter for the classification decision
- interested in knowing if there's any bias in the decision
- **Saliency Maps**, [Simonyan et al. (2013)]
 - computes the gradients of the predicted class scores with respect to the pixels of the input images
 - Gradients: for each pixel in the input images changes by a small amount, how much will the classification changes

Implementation

- Implemented using **TensorFlow** in Google Colab and trained via **GPU**
- Advantage of GPU: significantly reduced the training time
 - 10% of the data for 10 epochs via CPU: 30 hours
 - All data for 30 epochs via GPU: 20 minutes

More training details:

- During training, we used a **fixed number of epochs** as our stopping criterion
- Randomly shuffled the data after each epoch to create some randomness.
- The optimal set of parameters for each architecture was the set that has the minimum validation loss

Experiments and Results (Exploratory Analysis)



Figure: X-ray of a Normal Chest




Figure: X-ray of a Pneumonia Chest




Figure: Histogram and Contour




Figure: Histogram and Contour

Experiments and Results (Hyperparameters)

Table: Summary of Hyperparameters for each Neural Network

	Model 1	Model 2	Model 3
Batch size	64	64	64
# of epochs	30	30	30
Learning rate	10^{-3}	10^{-3}	10^{-3}
Filter size	NA	3×3	3×3
Stride (default)	NA	1	1
Dropout rate	NA	0.25	0.25

Experiments and Results (Final architecture)

The final architecture that was chosen

- **Model 1:**

- FLAT \rightarrow [90000]
- FC (Class scores)

Experiments and Results (Final architecture)

• Model 2:

- CONV1: 64 3×3 filters at stride 1, pad 0 $\rightarrow [298, 298, 64]$
- CONV2: 32 3×3 filters at stride 1, pad 0 $\rightarrow [296, 296, 32]$
- MAX POOL: 2×2 filter at stride 1 $\rightarrow [148, 148, 32]$
- FLAT $\rightarrow [700928]$
- FC (Class scores)




Figure: The Architecture for Model 2 (Simple CNN)

Experiments and Results (Final architecture)

● Model 3:

- CONV1: 64 3×3 filters at stride 1, pad 0 $\rightarrow [298, 298, 64]$
- NORM1 $\rightarrow [298, 298, 64]$
- AVE POOL $\rightarrow [149, 149, 64]$
- CONV2: 64 3×3 filters at stride 1, pad 0 $[147, 147, 64]$
- MAX POOL: 2×2 filter at stride 1 $[73, 73, 64]$
- FLAT $\rightarrow [341056]$
- FC $\rightarrow [128]$
- FC (Class scores)




Figure: The Architecture for Model 3 (Complex CNN)

Experiments and Results (Results)

Table: Experiment Results for each Neural Network

	Model 1	Model 2	Model 3
Total # of parameters	180,002	1,420,962	43,693,634
Size	2MB	16MB	500MB
Training Time	\approx 39 sec	\approx 10min	\approx 13min
Accuracy	71.47%	81.25%	81.09%
Precision	0.6906	0.7874	0.7688
Recall	0.9846	0.9590	0.9974

- Model 2 achieved the highest accuracy 81.25% among those 3 models
- Accuracy decreased slightly when increasing the complexity of the CNN models (from Model 2 to Model 3)
- recall (of Pneumonia) > 0.95 while precision < 0.8
(We care more on recall since we care among all the patients with Pneumonia, how many we detected v.s predicted some healthy people as Pneumonia)

Experiments and Results (Results)




Figure: One of the Pneumonia X-ray and the Corresponding Saliency Map

- The white areas in Figure 9 represents the large values of absolute gradients
- seems like the network reacted mostly to the areas of opacity in the chest and classified the X-ray as Pneumonia (as expected)

Conclusion

- Overall, Model 2 has the highest performance among those 3 models and is sensitive in detecting patients with Pneumonia.
- With high accuracy and high recall, model 2 can be served as initial screening to speed up the diagnosis process and let doctors to further investigate the identified Pneumonia cases
- Saliency maps didn't indicate decision making bias

References



He, K., Zhang, X., Ren, S., & Sun, J. (2015).

Delving deep into rectifiers: Surpassing human-level performance on imagenet classification.



Ioffe, S., & Szegedy, C. (2015)

Batch normalization: Accelerating deep network training by reducing internal covariate shift



Diederik P Kingma and Jimmy Ba. (2014)

Adam: A method for stochastic optimization



Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman (2013)

Deep inside convolutional networks: Visualising image classification models and saliency maps

The End